

CS100 Lecture 18

Contents

- Rvalue References
- Move Operations
 - Move Constructor
 - Move Assignment Operator
 - The Rule of Five
- `std::move`
- NRVO, Move and Copy Elision

Motivation: copy is slow

```
std::string a = some_value(), b = some_other_value();  
std::string s;  
s = a;  
s = a + b;
```

Consider the two assignments: `s = a` and `s = a + b`.

How is `s = a + b` evaluated?

Motivation: copy is slow

```
s = a + b;
```

1. Evaluate `a + b` and store the result in a temporary object, say `tmp`.
2. Perform the assignment `s = tmp`.
3. The temporary object `tmp` is no longer needed. Destroy it by calling its destructor.

Can we make this faster?

Motivation: copy is slow

```
s = a + b;
```

1. Evaluate `a + b` and store the result in a temporary object, say `tmp`.
2. Perform the assignment `s = tmp`.
3. The temporary object `tmp` is no longer needed. Destroy it by calling its destructor.

Can we make this faster?

- The assignment `s = tmp` is done by **copying** the contents of `tmp`.
- But `tmp` is about to die! Why can't we just *steal* the contents from it?

Motivation: copy is slow

Let's look at the other assignment:

```
s = a;
```

- **Copy** is needed here, because `a` lives long. It is not destroyed immediately after this statement is executed.
- You cannot just "steal" the contents from `a`. The contents of `a` must be preserved.

Distinguish between the different kinds of assignments

```
s = a;
```

```
s = a + b;
```

What is the key difference between them?

- `s = a` is an assignment from an **lvalue**,
- while `s = a + b` is an assignment from an **rvalue**.

If we only have the copy assignment operator, there is no way we can distinguish them.

* Define two different assignment operators, one accepting an lvalue and the other accepting an rvalue?

Rvalue References

A kind of reference that is bound to **rvalues**:

```
int &r = 42;           // Error: lvalue reference cannot be bound to rvalue
int &&rr = 42;          // Correct: `rr` is an rvalue reference
const int &cr = 42;    // Also correct:
                      // lvalue reference-to-const can be bound to rvalue
const int &&crr = 42;   // Correct, but useless
                      // rvalue reference-to-const is seldom used.

int i = 42;
int &&rr2 = i;          // Error: rvalue reference cannot be bound to lvalue
int &r2 = i * 42;       // Error: lvalue reference cannot be bound to rvalue
const int &cr2 = i * 42; // Correct
int &&rr3 = i * 42;     // Correct
```

- Lvalue references can only be bound to lvalues.
- Rvalue references can only be bound to rvalues.

Overload Resolution

```
void fun(const std::string &);  
void fun(std::string &&);
```

- `fun(s1 + s2)` matches `fun(std::string &&)`, because `s1 + s2` is an rvalue.
- `fun(s)` matches `fun(const std::string &)`, because `s` is an lvalue.
- Note that if `fun(std::string &&)` does not exist, `fun(s1 + s2)` also matches `fun(const std::string &)`.

Move Operations

Overview

The move constructor and the move assignment operator.

```
struct Widget {  
    Widget(Widget &&) noexcept;  
    Widget &operator=(Widget &&) noexcept;  
    // Compared to the copy constructor and the copy assignment operator:  
    Widget(const Widget &);  
    Widget &operator=(const Widget &);  
};
```

- Parameter type is **rvalue reference**, instead of lvalue reference-to-**const**.
- **noexcept** is necessary! (Will be covered in later lectures)

The Move Constructor

Take the `Dynarray` as an example.

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
public:
    Dynarray(const Dynarray &other) // copy constructor
        : m_storage(new int[other.m_length]), m_length(other.m_length) {
        for (std::size_t i = 0; i != m_length; ++i)
            m_storage[i] = other.m_storage[i];
    }
    Dynarray(Dynarray &&other) noexcept // move constructor
        : m_storage(other.m_storage), m_length(other.m_length) {
        other.m_storage = nullptr;
        other.m_length = 0;
    }
};
```

The Move Constructor

```
class Dynarray {  
    int *m_storage;  
    std::size_t m_length;  
public:  
    Dynarray(Dynarray &&other) noexcept // move constructor  
        : m_storage(other.m_storage), m_length(other.m_length) {  
  
    }  
};
```

1. *Steal* the resources of `other`, instead of making a copy.

The Move Constructor

```
class Dynarray {  
    int *m_storage;  
    std::size_t m_length;  
public:  
    Dynarray(Dynarray &&other) noexcept // move constructor  
        : m_storage(other.m_storage), m_length(other.m_length) {  
        other.m_storage = nullptr;  
        other.m_length = 0;  
    }  
};
```

1. *Steal* the resources of `other`, instead of making a copy.
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

* Take ownership of `other`'s resources!

The Move Assignment Operator

Take ownership of `other`'s resources!

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
  
        m_storage = other.m_storage; m_length = other.m_length;  
  
        return *this;  
    }  
};
```

1. *Steal* the resources from `other`.

The Move Assignment Operator

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
  
        m_storage = other.m_storage; m_length = other.m_length;  
        other.m_storage = nullptr; other.m_length = 0;  
  
        return *this;  
    }  
};
```

1. *Steal* the resources from `other`.
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

Are we done?

The Move Assignment Operator

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
  
        delete[] m_storage;  
        m_storage = other.m_storage; m_length = other.m_length;  
        other.m_storage = nullptr; other.m_length = 0;  
  
        return *this;  
    }  
};
```

0. Avoid memory leaks!

1. *Steal* the resources from `other` .
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

Are we done?

The Move Assignment Operator

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
        if (this != &other) {  
            delete[] m_storage;  
            m_storage = other.m_storage; m_length = other.m_length;  
            other.m_storage = nullptr; other.m_length = 0;  
        }  
        return *this;  
    }  
};
```

0. Avoid memory leaks!

1. *Steal* the resources from `other` .
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

* Self-assignment safe!

Lvalues are Copied; Rvalues are Moved

Before we move on, let's define a function for demonstration.

Suppose we have a function that concatenates two `Dynarray` s:

```
Dynarray concat(const Dynarray &a, const Dynarray &b) {  
    Dynarray result(a.size() + b.size());  
    for (std::size_t i = 0; i != a.size(); ++i)  
        result.at(i) = a.at(i);  
    for (std::size_t i = 0; i != b.size(); ++i)  
        result.at(a.size() + i) = b.at(i);  
    return result;  
}
```

Which assignment operator should be called?

```
a = concat(b, c);
```

Lvalues are Copied; Rvalues are Moved

Lvalues are copied; rvalues are moved ...

```
a = concat(b, c); // move assignment operator,  
                  // because concat(b, c) generates an rvalue.  
a = b; // copy assignment operator
```

Lvalues are Copied; Rvalues are Moved

Lvalues are copied; rvalues are moved ...

```
a = concat(b, c); // move assignment operator,  
                  // because concat(b, c) generates an rvalue.  
a = b; // copy assignment operator
```

... but rvalues are copied if there is no move operation.

```
// If Dynarray has no move assignment operator, this is a copy assignment.  
a = concat(b, c)
```

Synthesized Move Operations

Like copy operations, we can use `=default` to require a synthesized move operation that has the default behaviors.

```
struct X {  
    X(X &&) = default;  
    X &operator=(X &&) = default;  
};
```

- The synthesized move operations call the corresponding move operations of each member in the order in which they are declared.
- The synthesized move operations are `noexcept`.

The Rule of Five

The updated *copy control members*:

- copy constructor
- copy assignment operator
- move constructor
- move assignment operator
- destructor

If one of them has a user-provided version, the copy control of the class is thought of to have special behaviors.

The Rule of Five

- The move constructor or the move assignment operator will not be generated if any of the rest four members have a user-provided version.
- The copy constructor or copy assignment operator, if not provided by the user, will be implicitly `deleted` if the class has a user-provided move operation.
- The generation of the copy constructor or copy assignment operator is **deprecated** (since C++11) when the class has a user-provided copy operation or a destructor.

- This is why some of you see this error:

Implicitly-declared copy assignment operator is deprecated, because the class has a user-provided copy constructor.

The Rule of Five

The *copy control members* in modern C++:

- copy constructor
- copy assignment operator
- move constructor
- move assignment operator
- destructor

The Rule of Five: Define zero or five of them.

How to Invoke a Move Operation?

Suppose we give our `Dynarray` a label:

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
    std::string m_label;
};
```

The move assignment operator should invoke the **move assignment operator** on `m_label`. But how?

[illegible]

std::move

`std::move`

Defined in `<utility>`

`std::move(x)` performs an **lvalue to rvalue cast**:

```
int ival = 42;  
int &&rref = ival; // Error  
int &&rref2 = std::move(ival); // Correct
```

Calling `std::move(x)` tells the compiler that:

- `x` is an lvalue, but
- we want to treat `x` as an **rvalue**.

`std::move`

`std::move(x)` indicates that we want to treat `x` as an **rvalue**, which means that `x` will be *moved from*.

The call to `std::move` **promises** that we do not intend to use `x` again,

- except to assign to it or to destroy it.

A call to `std::move` is usually followed by a call to some function that moves the object, after which **we cannot make any assumptions about the value of the moved-from object**.

```
void foo(X &&x); // moves x
void foo(const X &x); // copies x
foo(std::move(x)); // matches foo(X&&), so that x is moved.
```

"`std::move` does not *move* anything. It just makes a *promise*."

Use `std::move`

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
    std::string m_label;
public:
    Dynarray(Dynarray &&other) noexcept
        : m_storage(other.m_storage), m_length(other.m_length),
          m_label(std::move(other.m_label)) {
        other.m_storage = nullptr;
        other.m_length = 0;
    }
};
```

Use `std::move`

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
    std::string m_label;
public:
    Dynarray &operator=(Dynarray &&other) noexcept {
        if (this != &other) {
            delete[] m_storage;
            m_storage = other.m_storage; m_length = other.m_length;
            m_label = std::move(other.m_label);
            other.m_storage = nullptr; other.m_length = 0;
        }
        return *this;
    }
};
```

Use `std::move`

Recall the **slow** string concatenation:

```
std::string s;  
for (int i = 0; i != n; ++i)  
    s = s + 'a';
```

- To compute `s + 'a'`, a **copy** of `s` is made because `s` is an **lvalue**. (Slow)
- The result of `s + 'a'` is stored in a **temporary object**, say `tmp`.
- Then the assignment `s = tmp` is performed, which is a **move** since `tmp` is an rvalue.

Use `std::move`

Recall the **slow** string concatenation:

```
std::string s;  
for (int i = 0; i != n; ++i)  
    s = s + 'a';
```

- To compute `s + 'a'`, a copy of `s` is made because `s` is an **lvalue**.
- Since the original value of `s` is not used anymore (after computing `s + 'a'`), is it possible to avoid this copy?

Use `std::move`

Require a move instead of copy, explicitly, by calling `std::move`.

```
std::string s;  
for (int i = 0; i != n; ++i)  
    s = std::move(s) + 'a';
```

- If the left-hand side object is an rvalue, this concatenation is directly performed by appending the right-hand side object to it.
 - This is reasonable, since the value of a *moved-from* object do not need to be preserved.

NRVO, Move and Copy Elision

Returning a Temporary (pure rvalue)

```
std::string foo(const std::string &a, const std::string &b) {  
    return a + b; // a temporary  
}  
std::string s = foo(a, b);
```

- First, a temporary is generated to store the result of `a + b`.
- How is this temporary returned?

Returning a Temporary (pure rvalue)

```
std::string foo(const std::string &a, const std::string &b) {  
    return a + b; // a temporary  
}  
std::string s = foo(a, b);
```

Since C++17, no copy or move is made here. The initialization of `s` is the same as

```
std::string s(a + b);
```

This is called **copy elision**.

Returning a Named Object

```
Dynarray concat(const Dynarray &a, const Dynarray &b) {  
    Dynarray result(a.size() + b.size());  
    for (std::size_t i = 0; i != a.size(); ++i)  
        result.at(i) = a.at(i);  
    for (std::size_t i = 0; i != b.size(); ++i)  
        result.at(a.size() + i) = b.at(i);  
    return result;  
}  
a = concat(b, c);
```

- `result` is a local object of `concat`.
- Since C++11, `return result` performs a **move initialization** of a temporary object, say `tmp`.
- Then a **move assignment** to `a` is performed.

Named Return Value Optimization, NRVO

```
Dynarray concat(const Dynarray &a, const Dynarray &b) {  
    Dynarray result(a.size() + b.size());  
    // ...  
    return result;  
}  
Dynarray a = concat(b, c); // Initialization
```

NRVO transforms this code to

```
// Pseudo C++ code.  
void concat(Dynarray &result, const Dynarray &a, const Dynarray &b) {  
    // Pseudo C++ code. For demonstration only.  
    result.Dynarray::Dynarray(a.size() + b.size()); // construct in-place  
    // ...  
}  
Dynarray a@; // Uninitialized.  
concat(a@, b, c);
```

so that no copy or move is needed.

Named Return Value Optimization, NRVO

Note:

- NRVO was invented decades ago (even before C++98).
- NRVO is an **optimization**, but not mandatory.
- Even if NRVO is performed, the move constructor should still be available. (Because the compiler can choose not to perform NRVO.)

Summary

In modern C++, unnecessary copies are greatly avoided by:

- move, with the returned lvalue treated as an rvalue, and
- NRVO, which constructs in-place the object to be initialized, and
- copy-elision, which avoids the move or copy of temporary objects.